


One node: $TP=4 \cdot DP\text{-}attention=4 \cdot PP=2$

Node 0 each shard holds $\frac{1}{4}$ of the requests + their KV cache


Stage 0 · layers 1–30 of 61 · $TP=4 \cdot DP\text{-}attn=4$



activations
NVLink



Stage 1 · layers 31–61 of 61 · $TP=4 \cdot DP\text{-}attn=4$



Node 1 · freed — serve a second replica

17,417 tok/s · 8 GPUs