



(a) TP=16 across nodes




11,253 tok/s · 16 GPUs

(b) TP=8 · PP=2 across nodes



14,632 tok/s · 16 GPUs (+30%)

(c) 1 node: TP=4 · DP=4 · PP=2 + DP-attn



Node 1 · freed — serve a second replica

17,417 tok/s · 8 GPUs (+55%)