

One node: $TP=4 \cdot DP\text{-}attention=4 \cdot PP=2$

Node 0 each shard holds $\frac{1}{4}$ of the requests + their KV cache

Stage 0 · layers 1-30 of 61 · $TP=4 \cdot DP\text{-}attn=4$

GPU 0

shard 1

Attn · full copy

GPU 1

shard 2

Attn · full copy

GPU 2

shard 3

Attn · full copy

GPU 3

shard 4

Attn · full copy

MoE experts · $TP=4 \cdot \frac{1}{4}$ per GPU



activations
NVLink



Stage 1 · layers 31-61 of 61 · $TP=4 \cdot DP\text{-}attn=4$

GPU 4

shard 1

Attn · full copy

GPU 5

shard 2

Attn · full copy

GPU 6

shard 3

Attn · full copy

GPU 7

shard 4

Attn · full copy

MoE experts · $TP=4 \cdot \frac{1}{4}$ per GPU

Node 1 · freed — serve a second replica

17,417 tok/s · 8 GPUs